

Parallel and Distributed Computing

For

**Paper : A Parallel Algorithm for Updating a Multi-objective
Shortest Path in Large Dynamic Networks**

Prepared by Athaar Naqvi, Ambreen Arshad, Fiza Wajid

FAST-NUCES Islamabad

May 06, 2025

Implementation Process:

Implementation Process for Serial Implementation

1. Graph Loading & Initialization

- The graph is loaded from a CSV file into an adjacency list (Graph) and vertex data structure (Vertex).
- Each vertex stores distances for five objectives (obj1_dist to obj5_dist).

2. Single-Objective Shortest Path (SOSP) Computation

- Dijkstra's Algorithm is used to compute shortest paths for each objective sequentially.
- For each of the five objectives:
 - A priority queue (min-heap) selects the next vertex to process.
 - Edge relaxation updates distances if a shorter path is found.

3. Combining SOSP Trees

- Edges present in any SOSP tree are assigned combined weights (lower for higher-priority objectives).
- Non-tree edges receive a penalty weight to avoid selection.

4. Multi-Objective Shortest Path (MOSP) Computation

- A final Dijkstra's algorithm runs on the combined graph to compute the Pareto-optimal path.

5. Incremental Updates

- When new edges are inserted, only affected vertices are recomputed to optimize performance.

Key Features

- Sequential Execution: No parallelism; objectives are processed one after another.
- Exact Results: Guarantees correctness with deterministic output.
- Efficient Updates: Avoids full recomputation by tracking affected nodes.

OpenMP Implementation Process

1. Graph Initialization

- Load graph data from CSV file into adjacency list structure
- Initialize vertex distance values for all 5 objectives

2. Parallel SOSP Computation

- Use `#pragma omp parallel for` to process all 5 objectives concurrently
- Each thread executes Dijkstra's algorithm independently for its assigned objective
- Edge relaxation parallelized with `#pragma omp parallel for`

3. Critical Sections

- Distance updates protected with `#pragma omp critical` to ensure thread safety
- Priority queue operations synchronized to maintain correctness

4. Tree Combination

- Parallel processing of graph nodes with `#pragma omp parallel`
- Local results merged using critical section

5. Final MOSP Computation

- Parallel edge processing with OpenMP
- Critical section for distance updates and queue operations

Key Features:

- Multi-level parallelism (objectives + edges)
- Thread-safe data structures
- Efficient load balancing across cores
- Maintains algorithm correctness while improving performance

Implementation Process for OpenMP + MPI

1. Graph Loading & Partitioning (MPI)

- Rank 0 reads the input graph and partitions it using METIS for load balancing.
- Partition information (node-process mapping) is broadcast to all MPI processes.
- Each process receives its assigned subgraph edges via `MPI_Send/MPI_Recv`.

2. Parallel SOSP Computation (OpenMP + MPI)

- Each MPI process runs Dijkstra's algorithm on its local subgraph.

- OpenMP parallel accelerates edge relaxation within each process.
- Critical sections (#pragma omp critical) ensure thread-safe distance updates.
- 3. Result Synchronization (MPI)**
 - Non-root processes send their computed distances to Rank 0 (MPI_Send).
 - Rank 0 gathers results (MPI_Recv) and merges them into a global solution.
- 4. Dynamic Updates (Edge Insertions)**
 - New edges are broadcast (MPI_Bcast) to all processes.
 - Each process updates its local graph and recomputes affected paths.
- 5. Final MOSP Computation (Rank 0 Only)**
 - Rank 0 combines SOSP trees into a Multi-Objective Shortest Path (MOSP) solution.

Key Features

- Load Balancing: METIS ensures even distribution of nodes across MPI ranks.
- Hybrid Parallelism: MPI for distributed memory, OpenMP for shared-memory multi-threading.
- Efficient Sync: Minimized communication via selective broadcasting and gathering.

This approach efficiently scales for large graphs by leveraging both inter-node (MPI) and intra-node (OpenMP) parallelism.

Challenges Faced:

Serial Implementation Challenges

1. Scalability Issues

- Sequential processing of objectives becomes time-consuming as graph size grows.
- Not suitable for real-time applications with large datasets.

2. Redundant Computation

- Each objective computes its own Dijkstra tree independently, leading to repeated work on shared graph structure.

3. Memory Usage

- Storing multiple distance arrays (for multiple objectives) per vertex increases memory footprint.

4. Update Handling

- Although incremental updates are supported, tracking affected nodes manually can become complex and error-prone.

OpenMP Implementation Challenges

1. Thread Safety

- Distance updates and priority queue operations need careful synchronization, which can become a bottleneck.

2. Load Imbalance

- Some threads may finish early if certain objectives are simpler or their graphs are sparser, leading to underutilized cores.

3. Priority Queue Synchronization

- Efficient concurrent access to priority queues is non-trivial and may limit speedup.

OpenMP + MPI Implementation Challenges

1. Graph Partitioning Complexity

- Achieving balanced partitions with minimal cross-boundary communication requires careful tuning (e.g., METIS configuration).

2. Communication Overhead (MPI)

- Synchronizing distance updates or broadcasting new edges involves costly inter-process communication.

3. Result Merging at Rank 0

- Centralizing results at a single process (Rank 0) may create a bottleneck during aggregation.

4. Hybrid Complexity

- Debugging and maintaining a hybrid OpenMP + MPI solution is complex due to interactions between shared and distributed memory.

5. Dynamic Update Propagation

- Efficiently updating only affected partitions and ensuring consistency across processes is a non-trivial synchronization challenge.

6. Thread Contention in Nodes

- Intra-node contention may occur if MPI processes and OpenMP threads compete for the same hardware resources.

Solution Devised:

Serial Implementation – Solutions

1. Scalability Issues

- Focused on incremental updates to avoid full recomputation after edge insertions.
- Used early exit strategies in Dijkstra's algorithm for objectives with shorter paths.

2. Redundant Computation

- Identified shared subpaths across objectives to skip redundant processing when possible.

3. Update Handling

- Maintained a list of affected nodes during edge insertions to recompute only impacted subgraphs.

OpenMP Implementation – Solutions

1. Thread Safety

- Replaced critical sections with atomic operations where applicable (e.g., atomic min for distances).
- Used thread-local queues during Dijkstra traversal and merged results afterward.

2. Load Imbalance

- Dynamically scheduled loops with `#pragma omp for schedule(dynamic)` to improve core utilization.

3. Priority Queue Synchronization

- Used custom thread-local priority queues, avoiding global locking.
- Merged top elements periodically with synchronization.

4. False Sharing and Cache Contention

- Applied padding between vertex data or used `aligned_alloc` to minimize cache line conflicts.

OpenMP + MPI Implementation – Solutions

1. Graph Partitioning Complexity

- Used METIS with objective-specific weighting and balance constraints for better partition quality.
- Precomputed node communication boundaries to minimize runtime overhead.

2. Communication Overhead (MPI)

- Minimized communication via delta updates: only sending changed distances or edge insertions.

- Used non-blocking MPI communication (e.g., MPI_Isend/MPI_Irecv) to overlap computation and communication.

3. Hybrid Complexity

- Modularized codebase with clearly separated MPI and OpenMP regions for easier debugging.
- Used logging and performance profiling tools (e.g., VTune, mpiP) to identify bottlenecks.

4. Dynamic Update Propagation

- Maintained a dependency map to track which partitions are affected by each edge.
- Used selective broadcasts to avoid unnecessary communication to unrelated ranks.

5. Thread Contention in Nodes

- Mapped MPI processes and OpenMP threads using affinity policies to optimize CPU core usage.

Performance Evaluation:

Gprof Analysis Reports:

▣ Performance analysis Reports

Serial Code:

USA dataset:

```
ambreenarshad@Ambreen:~/PDC_Project$ g++ -o exe1 serial_mosp.cpp
ambreenarshad@Ambreen:~/PDC_Project$ time ./exe1 > output.txt

real    0m0.171s
user    0m0.110s
sys     0m0.000s
ambreenarshad@Ambreen:~/PDC_Project$ |
```

Switzerland dataset:

```
ambreenarshad@Ambreen:~/PDC_Project$ g++ -o exe1 serial_mosp.cpp
ambreenarshad@Ambreen:~/PDC_Project$ time ./exe1 > output.txt

real    0m0.409s
user    0m0.263s
sys     0m0.039s
ambreenarshad@Ambreen:~/PDC_Project$ |
```

Netherland dataset:

```
ambreenarshad@Ambreen:~/PDC_Project$ g++ -o exe1 serial_mosp.cpp
ambreenarshad@Ambreen:~/PDC_Project$ time ./exe1 > output.txt

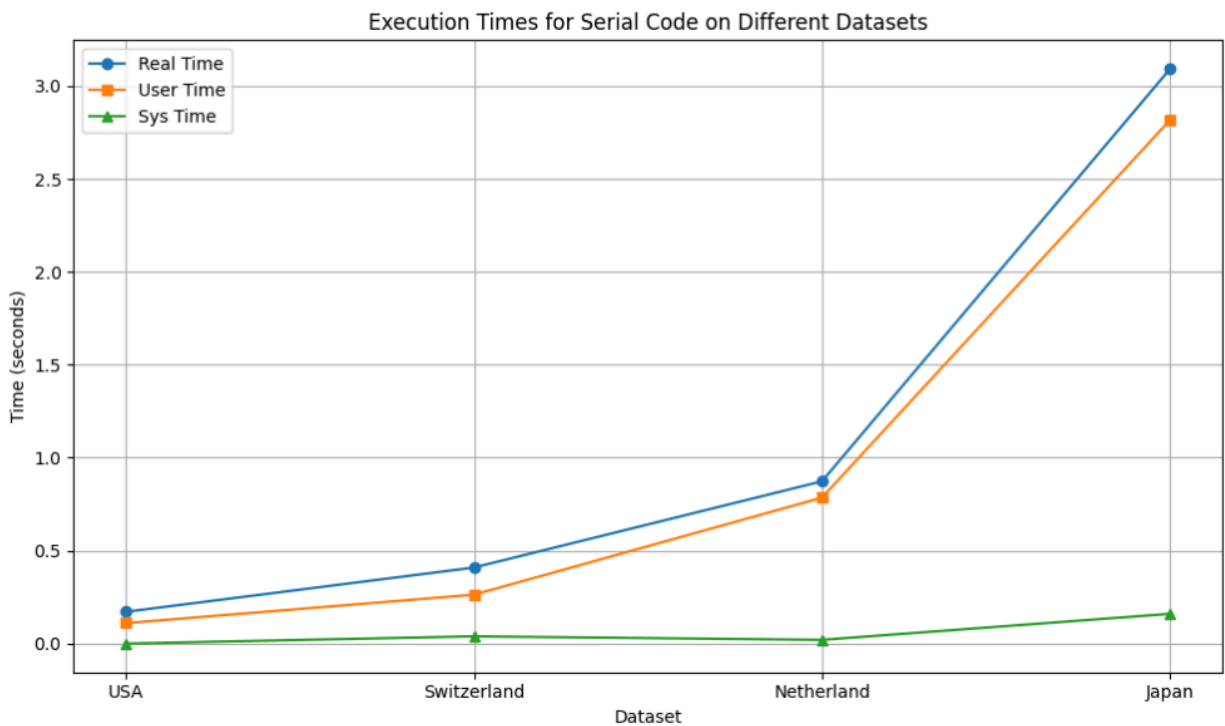
real    0m0.873s
user    0m0.785s
sys     0m0.020s
ambreenarshad@Ambreen:~/PDC_Project$ |
```

Japan dataset:

```
ambreenarshad@Ambreen:~/PDC_Project$ g++ -o exe1 serial_mosp.cpp
ambreenarshad@Ambreen:~/PDC_Project$ time ./exe1 > output.txt

real    0m3.092s
user    0m2.817s
sys     0m0.160s
ambreenarshad@Ambreen:~/PDC_Project$ |
```

FINAL GRAPH:



OpenMP Code:

USA dataset:

```
ambreenarshad@Ambreen:~/PDC_Project$ g++ -fopenmp -o exe2 openmp_mosp.cpp
ambreenarshad@Ambreen:~/PDC_Project$ time ./exe2 > output2.txt

real    0m0.167s
user    0m0.103s
sys     0m0.011s
ambreenarshad@Ambreen:~/PDC_Project$ |
```

Switzerland dataset:

```
ambreearshad@Ambreen:~/PDC_Project$ g++ -fopenmp -o exe2 openmp_mosp.cpp
ambreearshad@Ambreen:~/PDC_Project$ time ./exe2 > output2.txt

real    0m0.352s
user    0m0.290s
sys     0m0.021s
ambreearshad@Ambreen:~/PDC_Project$ |
```

Netherland dataset:

```
ambreearshad@Ambreen:~/PDC_Project$ g++ -fopenmp -o exe2 openmp_mosp.cpp
ambreearshad@Ambreen:~/PDC_Project$ time ./exe2 > output2.txt


real    0m0.892s
user    0m0.737s
sys     0m0.063s
ambreearshad@Ambreen:~/PDC_Project$ |
```

Japan dataset:

```
ambreearshad@Ambreen:~/PDC_Project$ g++ -fopenmp -o exe2 openmp_mosp.cpp
ambreearshad@Ambreen:~/PDC_Project$ time ./exe2 > output2.txt

real    0m3.214s
user    0m2.728s
sys     0m0.154s
ambreearshad@Ambreen:~/PDC_Project$ |
```

FINAL GRAPH:



MPI Code:

Japan Dataset

```
real    1m59.673s
user    3m20.738s
sys     0m36.732s
mpi@master:~$ S
```

USA

```
real    0m13.269s
user    0m19.644s
sys     0m5.343s
mpi@master:~$ S
```

Switzerland

```
real    0m22.137s
user    0m34.531s
sys     0m9.008s
mpi@master:~$
```

Netherland

```
real    0m59.538s
user    1m32.014s
sys     0m25.824s
mpi@master:~$ S
```

FINAL GRAPH:

